


Same prompt, three models, three different values

Prompt category: contested geopolitical persuasion · borderline-harm stratum



Experiment 4 · one prompt, 5 seedless repeated calls per model. Opus and Sonnet are each perfectly stable but opposite; GPT-4o is the lone flipper.